

Research Statement

Research agenda

My research asks how public institutions can use data-intensive systems without weakening accountability. I study the organisational decisions that surround a model: who defines the problem, what evidence counts as adequate, how errors are surfaced, and whether explanations change practice. This shifts attention from technical performance alone to the institutional conditions under which a system can be challenged and improved.

Current programme

My current work combines three methods. Audit studies reveal differences between stated governance procedures and operational decisions. Causal designs test whether specific forms of documentation or review alter those decisions. Participatory workshops show how affected communities define harms that are absent from administrative metrics. Together, these methods connect measurable effects with the reasons a policy succeeds or fails in practice.

One fictional project follows local-service teams as they review an automated eligibility tool. The study compares a standard model card with a structured public-audit protocol. The primary outcome is not whether reviewers report more information, but whether they identify consequential errors, change deployment conditions, and document unresolved disagreement.

Future work

The next phase will examine how oversight travels across organisations. I plan to compare procurement, internal review, and external audit as three pathways for institutional learning. The programme will produce open instruments, preregistered evaluations, and practical guidance co-developed with public organisations. Its broader contribution is a theory of accountable adoption: technical evidence matters, but its effects depend on who can act on it and what happens when evidence remains contested.